

A Hybrid Dynamical System for Vehicle Navigation

Comp 352

Micah Wylde

May 3, 2011

1 Introduction

Self-driving cars hold much promise for saving fuel, time and lives. Human-driven vehicles are responsible for millions of traffic accidents and over 30,000 deaths annually in the US alone. Humans can also make poor decisions in regards to routes, leading to traffic jams and wasted fuel. Safe, effective autonomous cars could largely solve these issues. AI researchers have been interested in the potential for vehicular navigation since the 1960s. In order to spur development in this area, the Defense Advanced Research Projects Agency (DARPA) organized the Grand Challenge in 2004, which tasked contestants to build cars which could autonomously navigate a 142 mile-long course in the Mojave desert. Though none of the vehicles could finish the course, a second competition the next year was more successful. Four teams finished in the allotted time, traversing a treacherous 132 mile desert course with no human guidance. Building on this achievement, in 2007 DARPA organized the Urban Challenge which took place in a simulated suburban environment. To win, cars had to navigate a maze of streets while accounting for other traffic and following California traffic laws at all times. Six teams finished the 61 mile course with the winner, “Boss” from CMU, taking a little more than four hours [1].

There are many challenges involved in autonomous driving and robot navigation in general. In the real world one must deal with unreliable and imperfect sensor data, imprecise localization technologies and other perception issues. Even in simulation the challenges of successful navigation are immense. Cars must follow traffic laws, get to their destination quickly and efficiently and act safely at all times, even in unexpected situations. As far as the actual navigation, there are two general approaches: deliberative and reactive. As an example of the former, A^* is an efficient and optimal graph search algorithm which is very effective at finding the best route between two points but cannot handle the dynamism of the real world. Dynamical systems-based navigation is a reactive strategy that uses force fields to guide agents away from obstacles and towards their target. However, it is purely local and has trouble finding optimal paths to distant goals. In this paper we present simulated agents which use each technique as well as a hybrid agent which makes use of A^* for global navigation and dynamical systems for local navigation.

1.1 Related works

Many group have created autonomous driving systems that can adapt to dynamic environments. Fraichard and Garnier [2] used fuzzy logic to adapt a pre-computed trajectory to current road conditions. They identify two components: a trajectory planner and an execution monitor. The trajectory planner is purely deliberative, although it works by a rule-based system rather than A^* . The EM is reactive and uses a base of fuzzy rules generated from human drivers. Like us, an action consists of an acceleration and a δ , which is the amount to turn the wheel. A more unusual approach to this problem is taken by Bell [3]. He describes an extension to A^* which instead of calculating the optimal path, finds the optimal *hyperpath*, which is a set

of paths any of which may be optimal, depending on future conditions. In particular, instead of considering only weights of edges (representing distances) we also consider an expected delay for each edge, which might represent traffic or other adverse road conditions. By pre-computing the hyperpath, we can adjust the route in real time as precise delay information becomes available.

More recently, “Boss” won the Urban Challenge using a sophisticated system which involves four systems: perception, motion planning, mission planning and behavioral execution. While the first and the last are mostly only relevant when dealing with real vehicles, the motion planner and mission planner are directly analogous to the two components of our hybrid agent. For the latter, which provides a series of way points along which the car should travel to get to its final destination, they use a process similar to ours A^* approach, albeit with a cost function that takes into account environmental complexity, expected time to traverse and other factors in addition to euclidean length of the edge. The motion planning system is responsible for choosing trajectories to the next way point. Their approach is not at all dynamical; it uses a model of vehicular motion to predict where the car (and other moving obstacles in the world) will be at some point in the future. The system chooses actions by optimizing parameters such that the differential equations that govern the car’s motion will evolve in the desired manner. In this sense their approach is almost deliberative but on very small time-scales. Additionally, after good parameters are chosen they are constantly modified as new information is fed into the system [4].

2 Methods

To create an autonomous agent that navigates while simulating a car’s behavior and traffic laws, we took three general approaches: deliberative planning through A^* search, reactive navigation through a dynamical system, and a hybrid of deliberative planning and reactive motion. The system architecture consists of a server and a separate client for each agent. Our goal was to produce agents which could navigate in a realistic environment while following traffic laws and acting safely at all times even when faced with dynamic, unexpected obstacles like pedestrians and other cars. Due to the great difficulty of producing the simulation, we were forced to scale back our ambitions to merely simulating driving in a relatively static world.

2.1 System Architecture

The project is written in Ruby, specifically JRuby¹ to utilize Java2D for the graphical display. As a result, the server code runs solely under JRuby, but client code can additionally be run using Ruby 1.9. Agents are represented both on the client and server end; server agents perform motion and display-related calculations, and client agents contain all the navigation inference and decision-making and ultimately send decisions (restricted to behavior variables) back to the server again. Agents communicate with the server over sockets by serializing hashes with YAML². Every part of the simulation is asynchronous: the display is constantly rendering the current state of the world as quickly as it can. The server agents are constantly recalculating their state variables according to their velocity and amount of turning (δ). The message loop is as follows: the client agents, which actually do the navigation computations, receive messages from the server with their current states. They then process that according to their operation and decide on an action, which is then sent to the server. The server carries out the action in the environment and sends the new state variables back to the agent.

The simulation is composed of many source files which control various pieces. The general architecture is given in figure 1, while the roles of the individual files are listed below. More information can be found in

¹JRuby is an implementation of the Ruby Programming Language on top of the Java Virtual Machine, which allows integration between Java and Ruby code

²YAML (a recursive acronym for YAML Ain’t Markup Language) is a language-independent data serialization standard which allows easy conversion of data structures to and from a string representation.

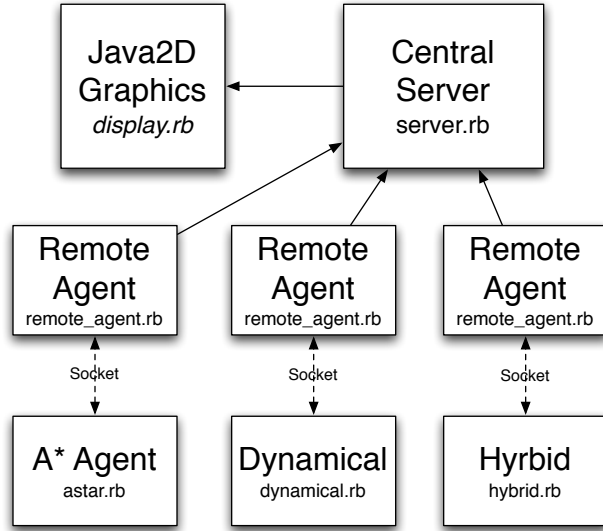


Figure 1: The system architecture of the simulation. Solid lines connect components included in the same process, while dotted lines indicated socket connections between discrete processes.

the documentation in each class.

app.rb The point of entry for the program. Processes command-line options and can optionally start a server, start a new agent, or run tests.

client_agent.rb Contains the `ClientAgent` class, the super class which every client agent inherits. Client agents receive messages from remote server agents, process the message (based on their form of navigation), and then send back a response in their behavior variable space.

constants.rb Contains various general constants that may be used in several locations or files, or that may be particularly useful to tweak with.

display.rb Contains all of the display code used to generate our rendering of the world.

map.rb Contains the classes, specifically `Map`, which encode information provided from real map data.

pqueue.rb A priority queue implementation used for A^* search.

remote_agent.rb Contains the `RemoteServerAgent` class, which is a subclass of `ServerAgent`. Essentially, a remote server agent is a server agent which is tied to a specific client agent and communicates with that client agent.

server_agent.rb Contains the `ServerAgent` class, which contains all the base representation and calculations needed for an agent (for example, the server agent computes various points needed to display the agent graphically).

server.rb Contains the socket server which handles agent connections.

socket.rb Contains code for handling communication between clients and the server

util.rb Contains a collection of geometry classes (`Point`, `Vector`, etc.) which are useful in the display and other various places (most notably in dynamical navigation calculations).

agents/astar.rb Contains a client agent that deliberately plans paths using A^* search.

agents/dynamical.rb Contains a client agent that navigates through a purely dynamical system-based approach.

agents/hybrid.rb Contains a client agent that navigates through a combination of deliberative planning and dynamical systems.

agents/simple.rb Contains a very primitive client agent (that can hardly be called an agent) which allows us to easily test the motion calculations performed by server agents.

2.2 Simulation



Figure 2: The Santa Cruz map displayed in the simulator, with an agent navigating on the left agent.

To provide a realistic environment for navigation tests, we used real-world map data from OpenStreetMap.org, which provides XML-formatted maps for the entire world. The data format, called OSM, provides a graph representation of a road system; streets are represented by placing nodes wherever the road turns or intersects other roads, with edges connecting each node. To provide testing data, we used the website to export maps of Hayward, CA and Santa Cruz, CA. We wrote a program, `osm_convert`, which converts an OSM XML file to a YAML format. In `map.rb` we then read this YAML data and construct an internal representation of the graph, converting points specified by latitude and longitude to a scale of meters.

In order to construct a simulation from this data, we needed to make roads from the edges of the graph. We did so by finding the normal vector to the edge and computing line segments that were a constant `ROAD_DIST` from the center-line. With the four corners, we could then draw the roads as polygons as well as use them for purposes of obstacle avoidance as described below. However, the naive implementation of this algorithm produced some overlapping walls and some walls which were too short. We therefore employ a clipping algorithm which ensures that the road edges meet smoothly. This does, however, leave empty spaces at intersections which appear green in the display.

The simulation has been designed for realism, with the agents responding in a physically plausible way according to two state variables which the agents can set: acceleration and δ , the amount that the wheel is turned. Raising δ above 0 causes the wheels to turn right and the car to follow in a circle, while lowering it below 0 causes the wheels to turn left. In this way we hope to demonstrate a direct analogy to control of robotic cars, which primarily concerns the acceleration/brake and steering. However we had difficulty

producing working navigation with this scheme, so the agents described below work by setting ϕ , the heading angle, directly.

2.3 Deliberative Agent




Figure 3: A^* navigation at work. The blue dots indicate the route.

We do deliberative navigation using the A^* graph search algorithm, which takes a heuristic function h and finds a path between two nodes on a graph. If h satisfies certain properties, the path returned will be optimal. We use the Euclidean distance formula $h(n_1, n_2) = \sqrt{(n_1.x - n_2.x)^2 + (n_1.y - n_2.y)^2}$ which satisfies these properties. The agent is modeled as a finite state machine with four states: start, straight, turn, and replan. The agent begins operation in the start state wherein it calculates the optimal route to its destination. Once that is done, it begins traveling forward until it is twice the road width from its starting node. It then transitions to straight mode, where it ensures that its ϕ is set parallel to the road. When it reaches a node, it pops the top node off the route stack and enters turn mode. In turn mode it sets its ϕ parallel to the vector between its target node (the node on the top of the route stack) and itself. This ensures the agent stays on the road at all times. Finally, when the agent has traveled past the node, it reverts back to straight mode. If the agent reaches a node which is not the next on its route, it enters replan mode, where it immediately stops and calculates a new route to its destination from its current position. Replanning can also be triggered by the server sending a new destination.

2.4 Reactive Agent

The reactive agent uses a dynamical systems approach for navigation. It forces the agent to stay in its lane by placing obstacles along the walls of the lane. We use a series of targets placed at the next intersection to guide the agent along the desired path. The most interesting detail of our dynamical system is our obstacle representation. The math of dynamical systems requires that all objects in the world be represented as circles, which is a poor fit for long straight obstacles such as walls. There are two conventional solutions to representing non-circular obstacles: use one large circle that bounds the obstacle or use many small circles along the obstacle. The first is untenable for the narrow corridors a car must pass through, while the second would require a very large number of circles and the corresponding computation. Our approach is to use a single circular obstacle per wall which follows the agent and which grows in size as the agent gets closer. Implementing this requires only basic geometry. We find the vector between the wall and the agent perpendicular to the road and displace the center of the circle from the road by a factor of the distance between the agent and the edge of the road.




Figure 4: Dynamical navigation at work. The red circles are obstacles that move along the road with the agent, while the blue circle is an intermediate target.

Like the deliberative agent, the reactive agent is modeled as a FSM with states start, normal, intersection and turn. Start waits until the agent is properly set up and chooses the initial target as the node the agent is currently facing, then transitions to normal mode. When in normal mode, the agent navigates until it reaches a node at which point it transitions to either intersection mode if the node it's reached has more than one neighbor or turn mode otherwise. If it's an intersection (which means more than one choice of which way to go) the agent computes the set $T = \{t \mid t \neq n_1, t \text{ is a neighbor of } n_0\}$ where n_0 is the node we've arrived at and n_1 is the node we've left. It then finds the $t \in T$ which minimizes $abs(u - v)$ where u is the vector from the agent to the destination and v is the angle from n_0 to t . We're trying to discover the best route using purely local information, so we find the path that appears to take us in the correct direction.

We set our dynamical parameters somewhat arbitrarily as follows: $m = 2, d_0 = 0.1, a = 20, \sigma = 1, h_1 = 10$. We do not use weights as we found them unnecessary. The most crucial parameter for our simulation seems to be a , the target attractor scaling factor. Lower a values do not produce working navigation.

2.5 Hybrid Agent

The hybrid agent was designed with the understanding that vehicular navigation is really a hybrid problem. Unlike more free-form environments, road systems are very limiting to dynamical navigation. Even though the agent may think its going in roughly the right direction, it has no way of knowing from purely local information whether it can even reach its destination by its path. Methods like A^* which look at global map data can solve this issue by finding the best path to take from every intersection so that the agent reaches the target in the minimum amount of time. However, deliberative methods do not handle the moment-to-moment navigation issues well, particularly in a dynamic environment. The best option seems to be a combination of the two: a hybrid dynamical agent. This agent behaves mostly like the dynamical agent described above except for how target selection works. Instead of choosing the next target using purely local information, we initially compute the A^* route as described in the deliberative section and use each node of that as the intermediate targets.

3 Results

We ran each agent 40 times on two different maps with random starting and destination points. Below are the percentage of runs that successfully found the target:

	Santa Cruz	Hayward
A*	24.1%	31.0%
Dynamical	37.2%	21.3%
Hybrid	36.2%	19.0%

These results are somewhat surprising, but likely speak more to bugs in the implementations of these strategies than the methods themselves. In particular, the very poor showing of hybrid navigation in the Hayward map is unexpected as is the general lack of success over all. From observation, the most common error conditions occur when a car accidentally goes off the road. At that point there is little the agents can do to rectify the situation, so once that happens there is little chance that the destination will be reached.

4 Conclusion

In the paper we have implemented three driving agents in a virtual city environment based on very different operating principles. A^* deliberates: at the beginning, when it receives its position and destination it does a length computation to calculate the optimal path to its goal, which it follows without interruption until something comes up and it must repeat the process. The dynamical agent, meanwhile, does no pre-computation and keeps little state between invocations of its sense/act function. It uses purely local information—what a driver could conceivably see out his or her windshield—to decide how to act at any given time. While the results above suggest that Dynamical is the best approach for getting to a target, that is likely the result of poor implementation. In fact, A^* knows where it must go at any intersection in order to obtain the shortest path while Dynamical must constantly make poor guesses from very limited information.

If we look past the problem of getting from our current position to our destination efficiently and consider the challenges of moment-to-moment navigation—how much to turn the wheel or press down the pedal at any given time—it’s clear that dynamical has the advantage. Not only can it react to changing map conditions (although this is not shown in the simulation) but from a design perspective it makes it much easier to get realistic motion. A^* relies on instantaneous change in the agent’s ϕ for its effectiveness. Forced to operate a wheel (i.e., setting δ instead of ϕ) the agent would have much greater difficulty figuring out how to turn at any given time in order to end up in the right place. Such control is much easier with Dynamical navigation, as it operates on $\dot{\phi}$, not ϕ (and $\phi \sim \delta$). Additionally we were unable to force the A^* agent to stay within its lane, while this was relatively simple with dynamical. In short, dynamical seems like a much better fit for the problem than a deliberative approach.

Dynamical’s weakness, as discussed above, is its global path planning which gives no guarantees of even reaching the destination. To rectify this, adding A^* ’s path planning to dynamical’s moment-to-moment navigation should produce an effective driving agent in our environment. And while this is the case in many situations, real-map data presents many edge conditions that can break it. Making all three agents more robust the issues in the map data is an important future direction, as is adding better recovery procedures. Currently the simulation is very static, and while many agents can be added (even running on different computers) they do not interact. Making a more vibrant city full of agents and other dynamic obstacles would present a much better comparison of the merits of each approach.

References

- [1] S. Thrun, “Toward robotic cars,” *Commun. ACM*, vol. 53, pp. 99–106, April 2010.
- [2] T. Fraichard and P. Garnier, “Fuzzy control to drive car-like vehicles,” *Robotics and Autonomous Systems*, vol. 34, no. 1, pp. 1 – 22, 2001.
- [3] M. G. Bell, “Hyperstar: A multi-path astar algorithm for risk averse vehicle navigation,” *Transportation Research Part B: Methodological*, vol. 43, no. 1, pp. 97 – 107, 2009.
- [4] C. Urmson, J. Anhalt, H. Bae, J. A. D. Bagnell, C. Baker, R. Bittner, T. Brown, M. N. Clark, M. Darms, D. Demitrish, J. Dolan, D. Duggins, D. Ferguson, T. Galatali, C. Geyer, M. Gittleman, S. Harbaugh, M. Hebert, T. Howard, S. Kolski, M. Likhachev, B. Litkouhi, A. Kelly, M. McNaughton, N. Miller, J. Nickolaou, K. Peterson, B. Pilnick, R. Rajkumar, P. Rybski, V. Sadekar, B. Salesky, Y.-W. Seo, S. Singh, J. Snider, J. Struble, A. Stentz, M. Taylor, W. R. Whittaker, Z. Wolkowicki, W. Zhang, and J. Ziglar, “Autonomous driving in urban environments: Boss and the urban challenge,” *Journal of Field Robotics Special Issue on the 2007 DARPA Urban Challenge, Part I*, vol. 25, pp. 425–466, June 2008.